


# IN4331 Web-scale Data Management

Group 7


Florentin Arsene  
Miruna Bețianu  
Pietro  
Mozafar Shah  
Pepijn Klop 4684311






# Tech Stack



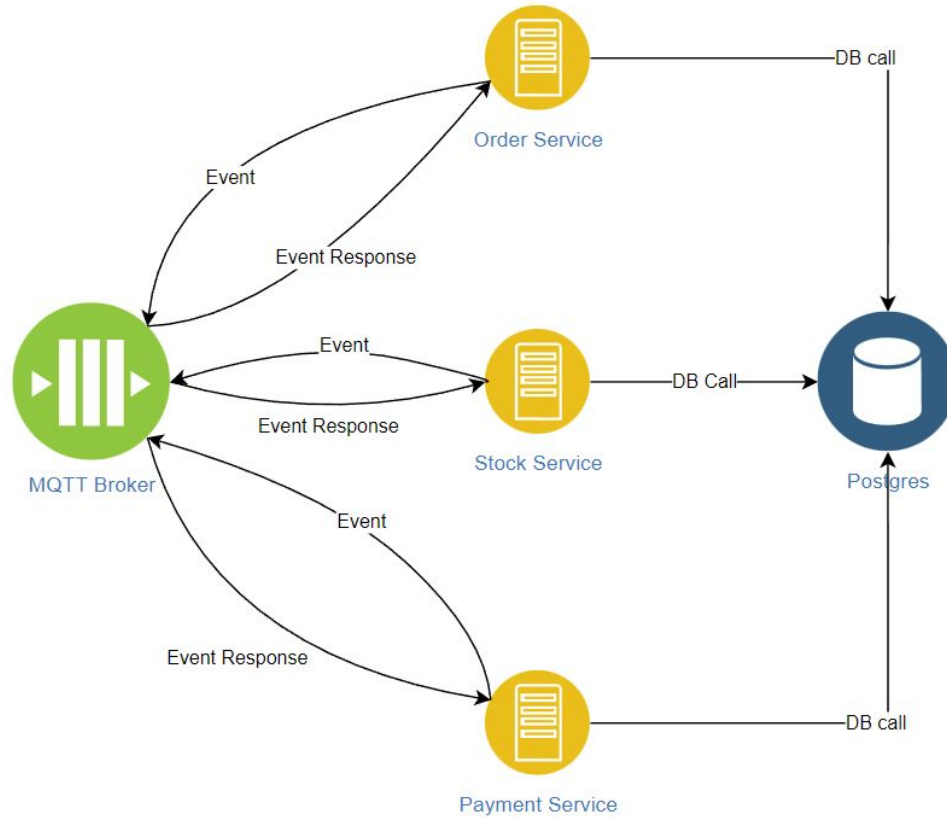
+ **Fiber** = ❤️



kubernetes





# Overall Design Architecture Diagram



# Orchestration VS Choreography

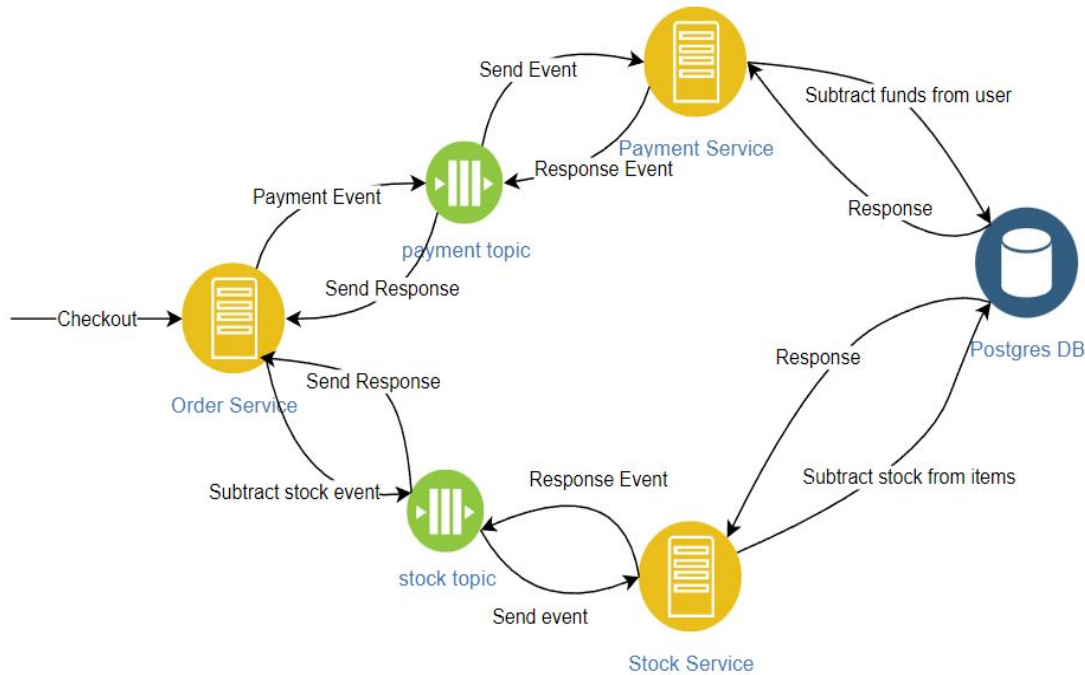


# Why we switched from HTTP calls



# Method of transaction execution


- Each transaction spans multiple services
- Transaction split up in local transactions
- Each service is responsible for guaranteeing local consistency
- Transaction is successful when all local transactions are successful.



# Guarantees offered by transactional design

- Compliance with ACID principles
- Consistency of data across the services
- Concurrency

How? Go routines!



## Method used for consistency

- Built-in methods:
  - PostgreSQL (synchronous commit)
- Implementation methods:
  - Message broker:
    - Messages
    - Publisher/Subscriber pattern
  - Single database



# Method of scalability

## Current implementation:

- Kubernetes
- Event based architecture

## Bottlenecks:

- Limited resources of local machine
- No load balancing
- Single database

# Method for fault-tolerance

## Database:

- Postgres fault-tolerant by design
- Transaction log for automatic crash recovery

## Service APIs:



- Error Detection
- Exception Handling

## Further improvements:

- Synchronous Replication
- Data block checksums
- Backup mechanisms

# Results

Max 200 RPS / 3600 users on a single CPU core  
Median response time of 8ms on every endpoint



# Conclusion

- Stable
- More than 200 RPS causes trouble
- Fault-tolerant
- Potential scalability and database improvements